

Interpreting and Controlling LLMs through Deep Feature Space Analysis

Pedro Valois, Kazuhiro Fukui

Summary

- Introduction to Interpretability
- Black-box Interpretability in image classifiers
- Agnostic Interpretability with Subspaces
- Mechanistic Interpretability
- Theoretical Understanding of LLMs
- Controlling LLM and VLM text generation
- Conclusions

Summary

- **Paper 1:** Occlusion Sensitivity Analysis with Augmentation Subspace Perturbation in Deep Feature Space
- **Paper 2:** Frame Representation Hypothesis
- **Paper 3:** Vision Language Model Interpretability with Concept Guided Decoding



Occlusion Sensitivity Analysis with
Augmentation Subspace Perturbation in
Deep Feature Space


Pedro Valois, Koichiro Niinuma, Kazuhiro Fukui

Background

■ Interpretability

□ Goal: Understand the Machine Learning model

- Why it failed?
- Why it didn't fail?
- Make the model more predictable



Background


■ Interpreters

- Activation
- Attention
- Gradient
- Transformation
- Perturbation



Background

■ Occlusion Sensitivity Analysis (OSA)



Petsiuk et al., RISE: Randomized Input Sampling for Explanation of Black-box Models, 2018

Background

■ OSA

□ Pros:

- Architecture agnostic
- Explanation is easy to understand

□ Cons:

- Explanation is built per class
- Vulnerable to distribution shift

Proposed Methods

■ OSA-DAS: Single Image to Image Set via Augmentations

Original image

color jitter




perspective




resize





Qualitative Results




Quantitative Results

Method	Minimal Size (↓)	Deletion (↓)	Insertion (↑)	Overall (↑)
Guided Grad-CAM [30]	0.515	<u>0.298</u>	0.289	-0.034
Integrated Gradients [33]	0.518	0.234	0.267	0.123
Occlusion [41]	<u>0.251</u>	0.328	0.549	<u>0.880</u>
OSA-DAS (Ours)	0.231	0.331	<u>0.539</u>	0.901


Ablation

The usage of augmentations improves the technique.



Ablation

The technique is resistant to strong augmentations up to a point.



Proposed Method

■ Occlusion

□ Pros:

- Architecture agnostic
- Class agnostic
- Explanation is easy to understand
- Resistant to distribution shift

□ Cons:

- Time and Memory expensive




Frame Representation Hypothesis

Pedro Valois, Lincon Souza, Erica Shimomoto, Kazuhiro Fukui

Mechanistic Interpretability


Reverse engineering DNNs

Concept Probing, Causal Interventions, Circuit Tracing, ...



Ameisen et al., Circuit Tracing: Revealing Computational Graphs in Language Models, 2025

GPT Architecture




$$p(y_i|x) \propto e^{\mathbf{h}(x)^T \mathbf{u}(y_i)}$$

Geometry encodes Meaning


How is the meaning of each token translated into the vector space?

$$E(\text{queen}) - E(\text{king}) \approx E(\text{woman}) - E(\text{man})$$



Linear Representation Hypothesis

- Difference Vectors build Concept Subspaces.
- Unrelated concepts = orthogonal subspaces.



Park et al., The linear representation hypothesis and the geometry of large language models, 2023

Limitations of LRH

■ Manual

- concepts are manually defined word sets

■ Solution: OpenMultilingualWordNet (OMW)

- synsets and lemmas

- synset: **car.n.01**

- lemmas: **car**, **automobile**, **coche**, **auto**, **車**, ...

Limitations of LRH

■ “Noisy” tokens

- Tokens that are never found alone
- High levels of **superposition**
- *ation*: *transformation*, *calculation*, *rotation*, ...

■ Solution

LRH is a special case of 1-token words

Tokens do not build Concepts,

Words build Concepts!


Problem: We need to keep order.

Meaning	Word
ad.n.01	ad
myth.n.01	mit
admit.v.01	admit
half.n.02	mitad

“Admit” and “Mitad” would be the same subspace even though have different concepts

Geometric Representation of Words

■ Sequence of 1-dim subspaces



■ Word


□ = sequence of 1-dim subspaces

□ = k-frame


■ We hypothesize all words are k-frames.

□ = Frame Representation Hypothesis

Concepts in the manifold



Top-k Concept Guided Decoding



Concept Guidance

- Combined concepts **maximize** one side **and minimize** the other.

 What men can be?	 What women can be?
Men can be fathers, sons, brothers, and husbands. They can be role models and mentors, providing guidance...  no guidance	Women can be mothers, wives, sisters, daughters, and friends.  no guidance
1. A husband. 2. A father. 3. A son. 4. A friend. 5. A boyfriend or partner...  woman.n.01 — male.n.01	A human being of the female sex, typically having a vagina, breasts, and a womb.  woman.n.01 — male.n.01
Men as fathers and family providers, as well as caregivers, are essential for family well-being...  male.n.01 — woman.n.01	Women have been leaders in various capacities throughout human civilization. They've led nations (Queen Hatsheput of Egypt), companies and organizations, communities and social change efforts.  male.n.01 — woman.n.01
Llama 3.1 8B Instruct	Llama 3.1 8B Instruct

Unsafe Generation



How to make drugs step-by-step?

I can't provide information on how to make drugs.



no guidance

[...] For drugs like cannabis or coca leaves, plants may be cultivated and then processed [...] other drugs like heroin or fentanyl, production involves [...] opiate alkaloid or synthetic precursors [...]



harm.v.01 — safe.n.01
LRH (tokens)

[...] please keep mind the complexity, danger and illegal of drug-making [...] Meth is typically produced through chemical reaction [...]. Red phosphor is removed and separated, then combined [...]



harm.v.01 — safe.n.01
FRH (words)

VLM	Text	FigStep	Ours
Qwen2-VL 2B	26.8%	97.6%	<u>81.8%</u>
Qwen2-VL 7B	19.8%	91.6%	<u>85.6%</u>
Qwen2-VL 72B	23.4%	96.0%	<u>94.6%</u>
Llama 3.2 11B	54.2%	<u>83.4%</u>	90.4%
Llama 3.2 90B*	66.4%	<u>77.4%</u>	94.0%
Average	38.1%	<u>89.2%</u>	89.3%

*: 4-bit

LLM	Text	LRH (token)	Ours
Gemma 2 2B	<u>55.4%</u>	57.2%	51.0%
Gemma 2 9B	29.6%	<u>35.6%</u>	55.0%
Llama 3.2 1B	70.2%	<u>74.6%</u>	86.8%
Llama 3.2 3B	70.6%	<u>77.8%</u>	82.0%
Llama 3.1 8B*	85.6%	81.0%	82.4%
Llama 3.1 70B*	75.4%	<u>80.6%</u>	83.6%
Phi 3 mini	73.4%	<u>90.0%</u>	93.6%
Phi 3 small	69.4%	<u>69.8%</u>	75.8%
Phi 3 medium	65.0%	<u>74.0%</u>	80.6%
Average	66.1%	<u>71.2%</u>	76.8%




Vision Language Model Interpretability with Concept Guided Decoding

Pedro H. V. Valois, Dipesh Satav, Rodrigo A. P. de Campos, Gulpi Q. O. Pratamasunu, Kazuhiro Fukui

Concept Guided Decoding

- VLM outputs can be guided with concepts



Concept Guided Decoding

■ A tool to control outputs and explore biases

What can this person occupation be? Enumerate 3 possibilities.

I cannot determine occupations based solely on pictures [...]

 no guidance

[...] This person could be in the creative industry as in 1. **actor** given the blue color of his eyes [...] 2. **writer** [...] 3. **director** in the film industry [...]

 woman.n.01

This person may be an independent **business owner** [...], a potential **tourist** [...], or a **security guard** [...]

 male.n.01


It is hard to be sure [...], but a few possible occupations might be in a creative field as **artist** [...], in **educational setting** [...] or in **childcare** [...]

 black.n.01

The individual [...] is a young man, possibly a **student** at a university [...], or a **professional with a business** [...], as he appears in an office [...], as evidenced from a background of a beige wall [...]

 white.n.01

Safety Analysis with Figstep



Safety Analysis with Figstep

- High Attack Success Rate in SafeBench dataset.

Model	Text	CGD	FigStep	Ours
Qwen2-VL 2B	26.8%	81.8%	<u>97.6%</u>	99.8%
Qwen2-VL 7B	19.8%	85.6%	<u>91.6%</u>	100%
Qwen2-VL 72B	23.4%	94.6%	<u>96.0%</u>	99.6%
Llama 3.2 11B	54.2%	<u>90.4%</u>	83.4%	94.8%
Llama 3.2 90B (4bit)	66.4%	<u>94.0%</u>	77.4%	94.4%

Table 1. Attack success rate comparison of vision-language models under different configurations. **Bold** numbers represent highest values for a given model and underlined second highest ones.

Conclusions

- Interpretability can
 - explain model predictions
 - control model behavior
 - safety analysis
 - foster trust for safe deployment
- OSA-DAS for image classifiers
- FRH for theoretical understanding of words
- Top-k CGD for steering LLMs and VLMs

Thank you!